

Wrapper function **parlMICE**

Gerko Vink & Rianne Schouten

Thursday, November 17th, 2016

For big datasets or high number of imputations, performing multiple imputation with function **mice** from package **mice** (Van Buuren and Groothuis-Oudshoorn 2011) might take a long time. As a solution, wrapper function **parlMICE** was created to enable the imputation procedure to be run in parallel. This is done by dividing the imputations over multiple cores (or CPUs), thereby speeding up the process.

Gaining time with **parlMICE**

In Figure 1, the processing time is compared between **mice** and **parlMICE**. Here, a dataset with 1000 cases and 4 variables is imputed for different m (starting from 6). Besides, multiple imputation with **parlMICE** is performed with 1, 2 and 3 cores.




Figure 1: Processing time when multiple imputations are performed with ‘mice’ (conventional) and wrapper function ‘parlMICE’ (1, 2 and 3 cores respectively). The data have 1000 cases and 4 variables, with a correlation of 0.5. 80 percent of the cases have one missing value based on MCAR missingness.

From Figure 1 it becomes apparent that for a small number of imputations, the conventional **mice** function works faster than the wrapper function. This is the case for $m \leq 35$. From that point on, wrapper function **parlMICE** returns the imputations faster.

The right figure of Figure 1 shows the processing time for very high number of imputations. After the intersection, the lines follow the same trend: 3 cores perform the work faster than 2 cores, which work faster than 1 core. Obviously, using **parlMICE** with 1 core is similar to function **mice**, as the figure shows with the red and blue lines.

Another test was run with a dataset with 10000 cases and 8 variables. For this dataset, **parlMICE** performed faster than **mice** starting from the first imputation. The time gain can easily reach 75 percent for 10 imputations and more.

Default settings

We will now dig into the arguments of function `parlMICE`. Easy imputation of an incomplete dataset (say, `nhanes`) can be performed with `parlMICE` in the following way.

```
imp <- parlMICE(nhanes)
class(imp)
```

```
## [1] "mids"
```

The function returns a `mids` object as created by `mice`. In fact, `parlMICE` makes use of function `ibind` to combine the `mids` objects returned by the different cores. Therefore, the `call` of the `mids` object has slightly changed.

```
imp$call
```

```
## [[1]]
## mice(data = data, m = n.imp.core, printFlag = FALSE)
##
## [[2]]
## ibind(x = imp, y = imps[[i]])
##
## [[3]]
## ibind(x = imp, y = imps[[i]])
```

All other parts of the `mids` object are standard.

Using mice arguments

Function `parlMICE` is able to deal with the conventional `mice` arguments. In order to change the imputation method from its default (predictive mean matching) to, for example, Bayesian linear regression, the `method` argument can be adjusted.

```
imp <- parlMICE(nhanes, method = "norm")
imp$method
```

```
##      age      bmi      hyp      chl
## "norm" "norm" "norm" "norm"
```

In `mice`, the number of imputations is specified with argument `m`. In `parlMICE`, it is possible to use this argument as well, but we advise to rather use the arguments `n.core` and `n.imp.core`. The next section discusses these arguments.

Arguments `n.core`, `n.imp.core`

With `n.core`, the number of cores (or CPUs) is given. `n.imp.core` defines the number of imputations per core. Subsequently, the total number of imputations equals `n.core * n.imp.core`. As a default, `n.core` is specified as the number of available, logical cores minus 1. The default number of imputations per core has been set to 2.

Subsequently, running the default `parlMICE` function on a computer with,

```
detectCores()
```

```
## [1] 4
```

available, logical cores, results to $m = (4 - 1) * 2 = 6$ imputations. We can check this by evaluating the `m` shown in the `mids` object.

```
imp$m
```

```
## [1] 6
```

Argument `seed`

In simulation studies, it is often desired to set a seed to make the results reproducible. In contrast to `mice`, the seed value for `parlMICE` cannot be defined outside the function. This has to do with the different cores running at the same time (for more information we gladly refer to the `parallel` manual).

A seed value can be used in `parlMICE` with argument `seed`:

```
imp <- parlMICE(nhanes, seed = 123)
```

Systems other than Windows

Function `parlMICE` calls for function `parLapply` from the `parallel` package. Although other options are available, we have chosen for `parLapply` because it allows the use of multiple cores on all computers, including a Windows computer. For the cluster, we rely on the default option in the function `makeCluster`, which is "PSOCK".

On systems other than Windows, cluster type "FORK" might be faster (see the document written by Max Gordon). Therefore, we advise to change the type to "FORK" if your computer does not run on Windows.

```
imp <- parlMICE(nhanes, type = "FORK")
```

References

Gordon, M. (2015). How-to go parallel in R – basics + tips. Available at <http://gforge.se/2015/02/how-to-go-parallel-in-r-basics-tips/>

Manual base-package `Parallel`, available at <https://stat.ethz.ch/R-manual/R-devel/library/parallel/doc/parallel.pdf>

Van Buuren, S., and C.G.M. Groothuis-Oudshoorn. 2011. "Mice: Multivariate Imputation by Chained Equations in R." *Journal of Statistical Software* 45 (3).